

Assignment2

May 28, 2022

Perform Comparative evaluation of Naïve Bayes, Decision Tree, Random Forest, and Support Vector Machine classifier. Use diabetes and hear disease datasets

0.1 Loading the Required Libraries

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import classification_report
from sklearn.metrics import plot_confusion_matrix
from sklearn import model_selection
import matplotlib.pyplot as plt
import warnings

warnings.filterwarnings('ignore')
%matplotlib inline
```

1 Using Diabetes Dataset

Reading the CSV File as a Pandas DataFrame

```
[2]: df = pd.read_csv('data/diabetes.csv')
```

Checking the first few rows

```
[3]: df.head()
```

```
[3]:   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin   BMI  \
0             6     148             72             35         0  33.6
1             1      85             66             29         0  26.6
2             8     183             64              0         0  23.3
3             1      89             66             23        94  28.1
4             0     137             40             35       168  43.1
```

DiabetesPedigreeFunction Age Outcome

0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

Plotting the Bar Graph of Target Class

```
[4]: df['Outcome'].value_counts().plot(kind='bar')
plt.title('Bar Plot of Target Class')
```

```
[4]: Text(0.5, 1.0, 'Bar Plot of Target Class')
```



Checking if there is null values in the dataset

```
[5]: df.isna().sum()
```

```
[5]: Pregnancies      0
      Glucose          0
      BloodPressure    0
      SkinThickness     0
      Insulin           0
      BMI              0
      DiabetesPedigreeFunction  0
      Age              0
```

```
Outcome          0
dtype: int64
```

There are no an null values in the dataset

Train Test Split

```
[6]: X = df.drop(columns=['Outcome'],axis=1)
     y = df['Outcome']
```

```
[7]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
     ↪2,random_state=0)
```

A helper function that runs all the machine models and gives the evaluation matrices as output

```
[8]: def run_exps(X_train, y_train, X_test, y_test):
     dfs = []
     # Defining the Models and adding them to a list
     models = [
         ('NaiveBayes',GaussianNB()),
         ('DecisionTree',DecisionTreeClassifier()),
         ('RandomForest', RandomForestClassifier()),
         ('SVM', SVC())

     ]
     results = []
     names = []
     scoring = ['accuracy', 'precision', 'recall', 'f1']
     target_names = ['0','1']
     # Looping through Each Model in the list and making prediction to get
     ↪evaluation matrices
     for name, model in models:
         kfold = model_selection.KFold(n_splits=5, shuffle=True, random_state=0)
         cv_results = model_selection.cross_validate(model, X_train, y_train,
         ↪cv=kfold, scoring=scoring)
         clf = model.fit(X_train, y_train)
         y_pred = clf.predict(X_test)
         print("#####")
         print(name)
         print(classification_report(y_test, y_pred, target_names=target_names))
         plot_confusion_matrix(clf, X_test, y_test)
         plt.title(name)
         results.append(cv_results)
         names.append(name)
         this_df = pd.DataFrame(cv_results)
         this_df['model'] = name
         dfs.append(this_df)
         final = pd.concat(dfs, ignore_index=True)
```

```
return final
```

```
[9]: r_df = run_exps(X_train,y_train,X_test,y_test)
```

```
#####
```

```
NaiveBayes
```

	precision	recall	f1-score	support
0	0.84	0.87	0.85	107
1	0.67	0.62	0.64	47
accuracy			0.79	154
macro avg	0.76	0.74	0.75	154
weighted avg	0.79	0.79	0.79	154

```
#####
```

```
DecisionTree
```

	precision	recall	f1-score	support
0	0.86	0.79	0.83	107
1	0.60	0.70	0.65	47
accuracy			0.77	154
macro avg	0.73	0.75	0.74	154
weighted avg	0.78	0.77	0.77	154

```
#####
```



```
RandomForest
```



	precision	recall	f1-score	support
0	0.83	0.89	0.86	107
1	0.70	0.60	0.64	47
accuracy			0.80	154
macro avg	0.77	0.74	0.75	154
weighted avg	0.79	0.80	0.79	154

```
#####
```

```
SVM
```

	precision	recall	f1-score	support
0	0.81	0.92	0.86	107
1	0.73	0.51	0.60	47
accuracy			0.79	154
macro avg	0.77	0.71	0.73	154
weighted avg	0.78	0.79	0.78	154





Getting the Average of Each Matrices from K-Fold CV of each Model

```
[10]: r_df.groupby('model').mean().reset_index()
```

```
[10]:
```

	model	fit_time	score_time	test_accuracy	test_precision	\
0	DecisionTree	0.004315	0.004933	0.698707	0.585716	
1	NaiveBayes	0.004870	0.007758	0.747621	0.669703	
2	RandomForest	0.186305	0.019301	0.750846	0.687880	
3	SVM	0.011196	0.008482	0.741090	0.710155	

	test_recall	test_f1
0	0.558024	0.565386
1	0.600590	0.632148
2	0.568713	0.621267
3	0.470610	0.564164

From the above results we can say that RandomForest is the best model on the basis of accuracy.

2 Using Heart Disease Dataset

Reading the CSV File as a Pandas DataFrame

```
[11]: df = pd.read_csv('data/heart.csv')
```

Checking First few rows of the dataset

```
[12]: df.head()
```

```
[12]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	52	1	0	125	212	0	1	168	0	1.0	2	
1	53	1	0	140	203	1	0	155	1	3.1	0	
2	70	1	0	145	174	0	1	125	1	2.6	0	
3	61	1	0	148	203	0	1	161	0	0.0	2	
4	62	0	0	138	294	1	1	106	0	1.9	1	

	ca	thal	target
0	2	3	0
1	0	3	0
2	0	3	0
3	1	3	0
4	3	2	0

Converting the sex feature to categorical

```
[13]: df['sex'] = df['sex'].astype('category')
```


Checking for missing values in the dataset

```
[14]: df.isna().sum()
```

```
[14]: age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

There is no missing values in the dataset

```
[15]: df['target'].value_counts().plot(kind='bar')
plt.title('Bar Graph of Target Values')
plt.show()
```



Train Test Split


```
[16]: X = df.drop(columns=['target'])
      y = df['target']
```

```
[17]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
      ↪3,random_state=0)
```

Calling the Helper Function

```
[18]: r_df = run_exps(X_train,y_train,X_test,y_test)
```

```
#####
```

NaiveBayes

	precision	recall	f1-score	support
0	0.85	0.81	0.83	145
1	0.84	0.87	0.86	163
accuracy			0.84	308
macro avg	0.84	0.84	0.84	308
weighted avg	0.84	0.84	0.84	308

```
#####
```

DecisionTree

	precision	recall	f1-score	support
0	1.00	1.00	1.00	145
1	1.00	1.00	1.00	163
accuracy			1.00	308
macro avg	1.00	1.00	1.00	308
weighted avg	1.00	1.00	1.00	308

```
#####
```

RandomForest


	precision	recall	f1-score	support
0	0.98	1.00	0.99	145
1	1.00	0.98	0.99	163
accuracy			0.99	308
macro avg	0.99	0.99	0.99	308
weighted avg	0.99	0.99	0.99	308



```
#####
```


SVM

	precision	recall	f1-score	support
0	0.75	0.70	0.73	145

	1	0.75	0.79	0.77	163
accuracy				0.75	308
macro avg		0.75	0.75	0.75	308
weighted avg		0.75	0.75	0.75	308







Getting Average score from K-Fold CV result for each model

```
[19]: r_df.groupby('model').mean().reset_index()
```

```
[19]:
```

	model	fit_time	score_time	test_accuracy	test_precision \
0	DecisionTree	0.003666	0.004892	0.972067	0.976604
1	NaiveBayes	0.004914	0.006603	0.811694	0.796276
2	RandomForest	0.193031	0.021173	0.969279	0.973574
3	SVM	0.019332	0.011124	0.694610	0.685327

	test_recall	test_f1
0	0.965967	0.971049
1	0.842783	0.818564
2	0.962835	0.967949
3	0.742913	0.710625

From the result above we can say that Random Forest is the best model as the values of accuracy, precision and recall is the highest.